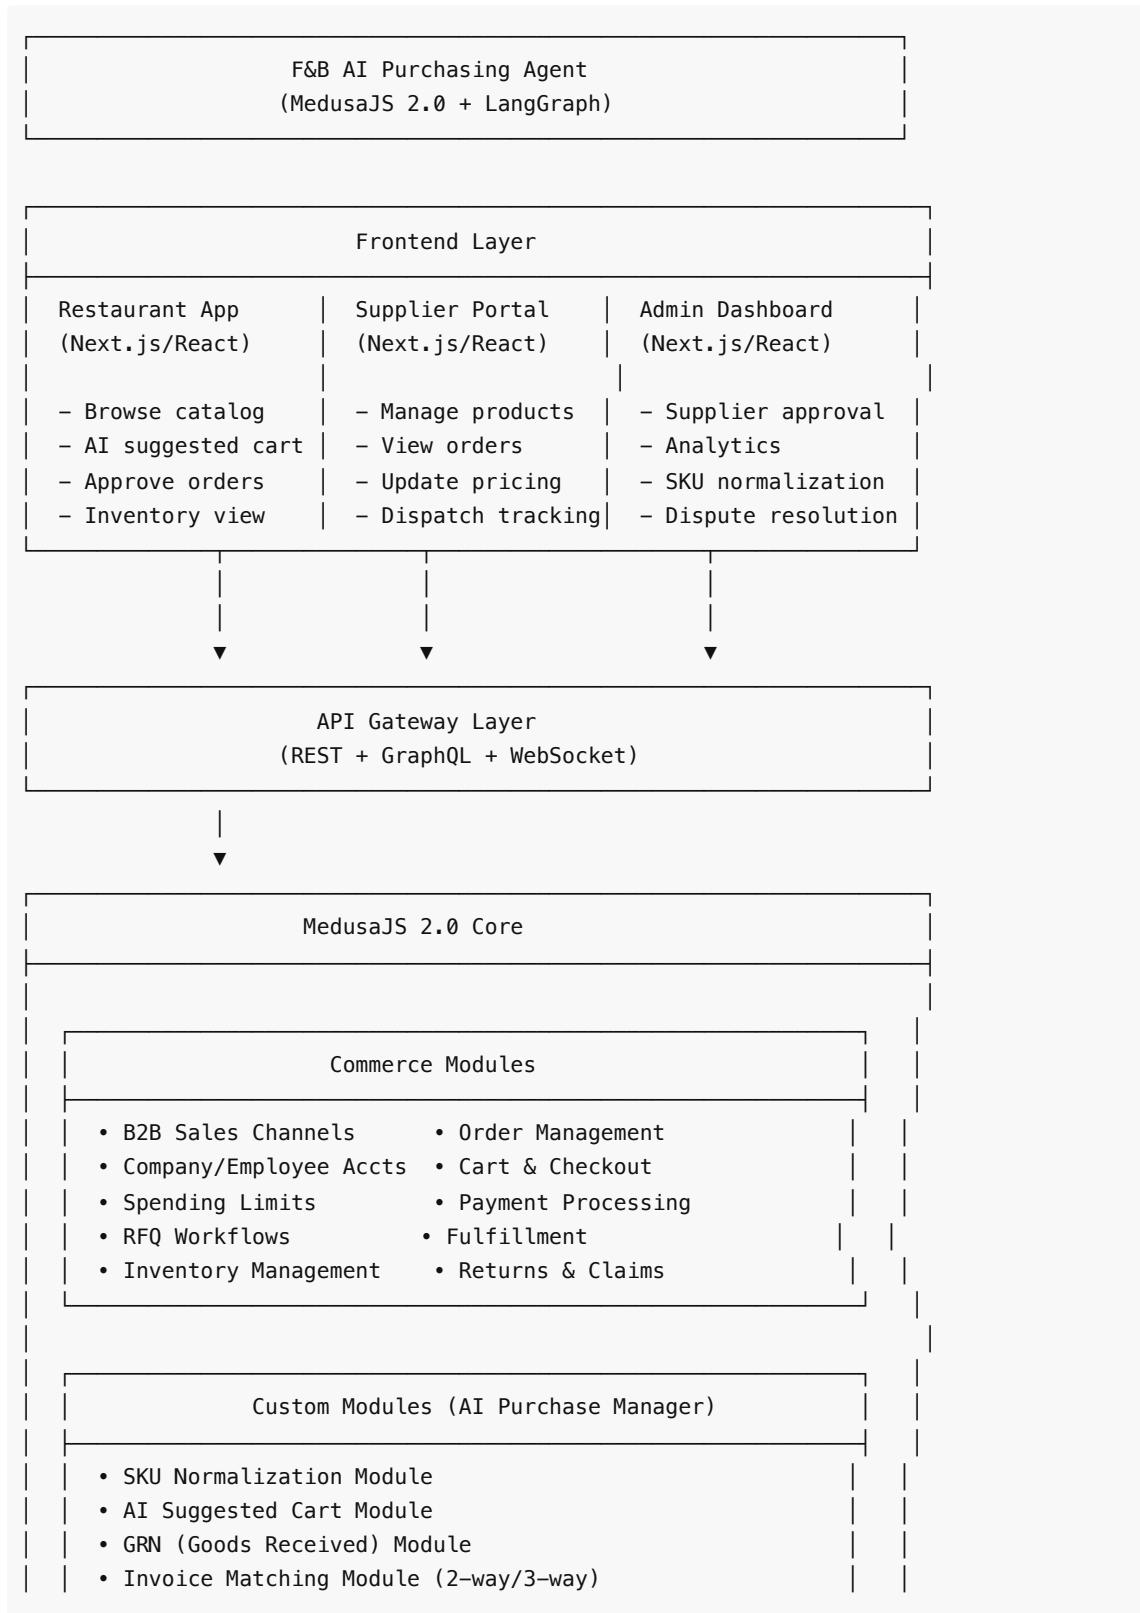
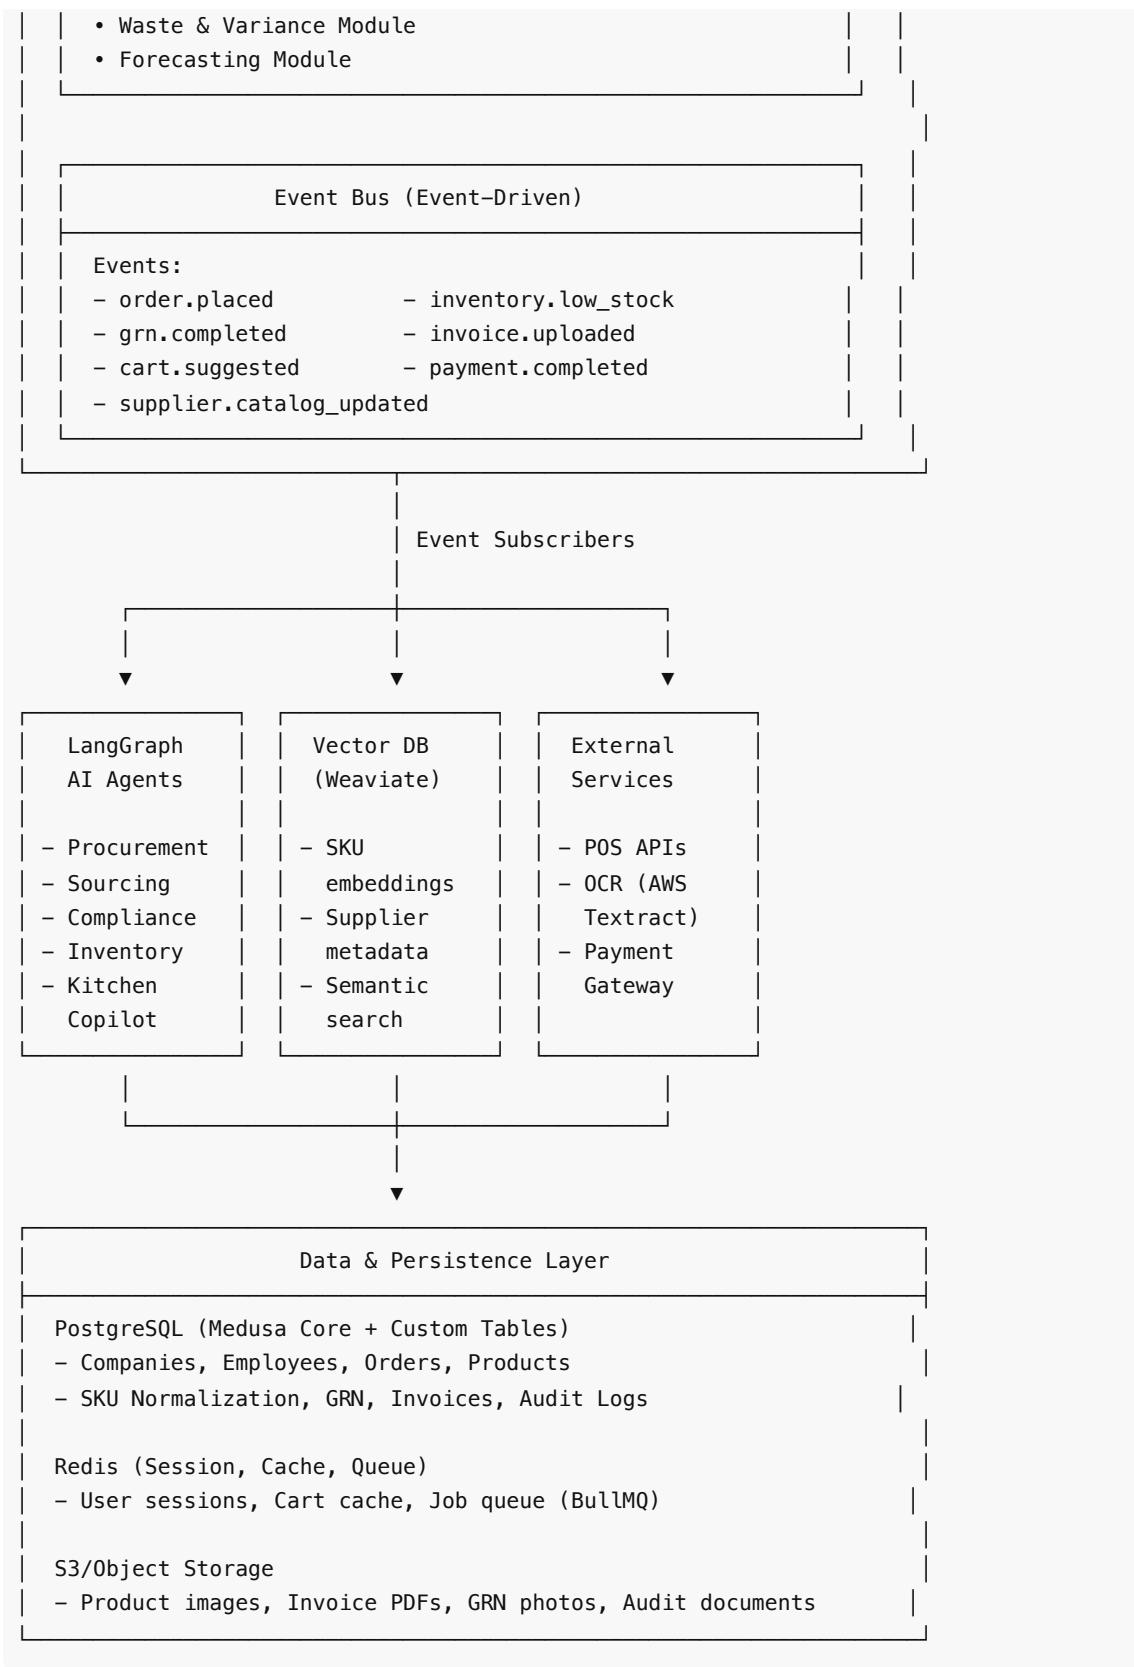


MedusaJS 2.0 + LangGraph Architecture

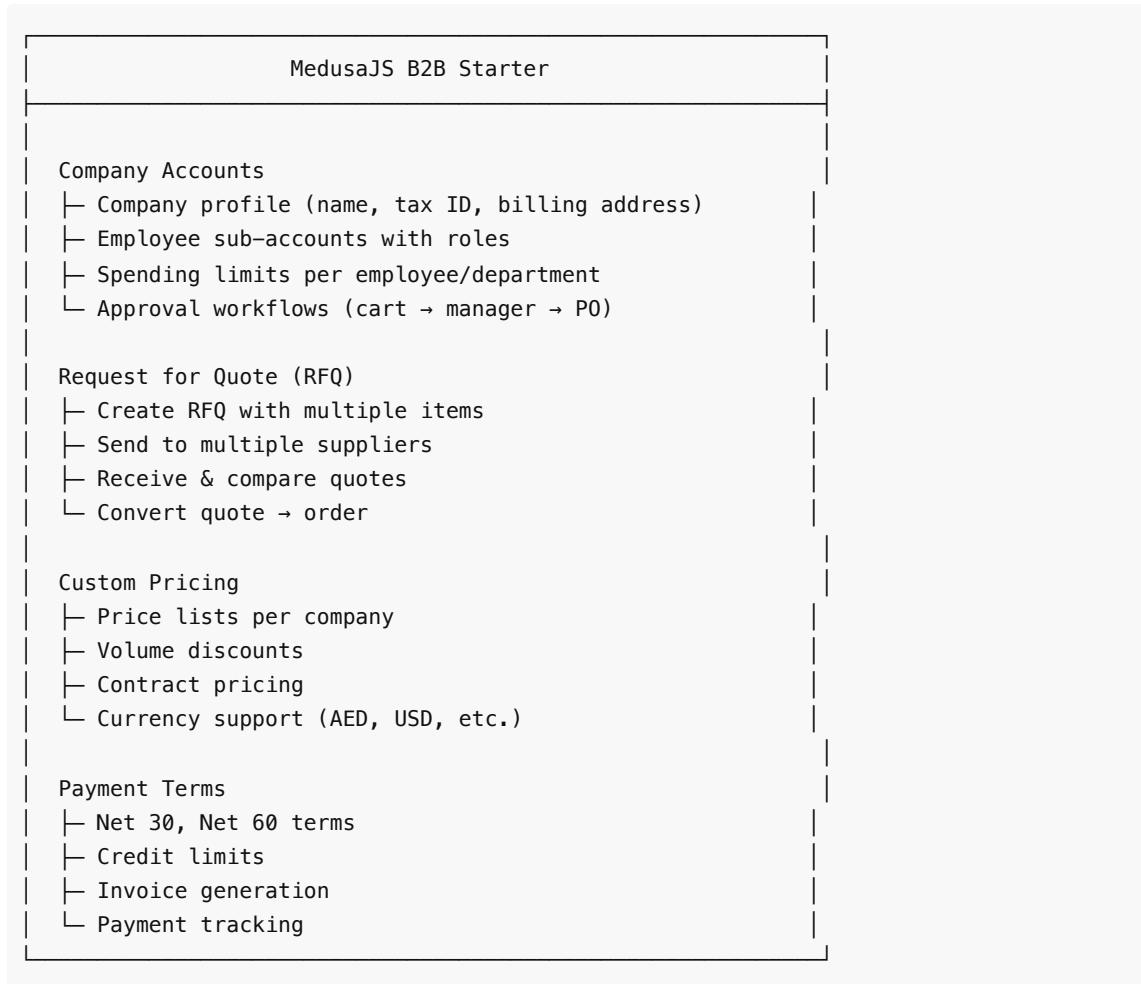
System Overview





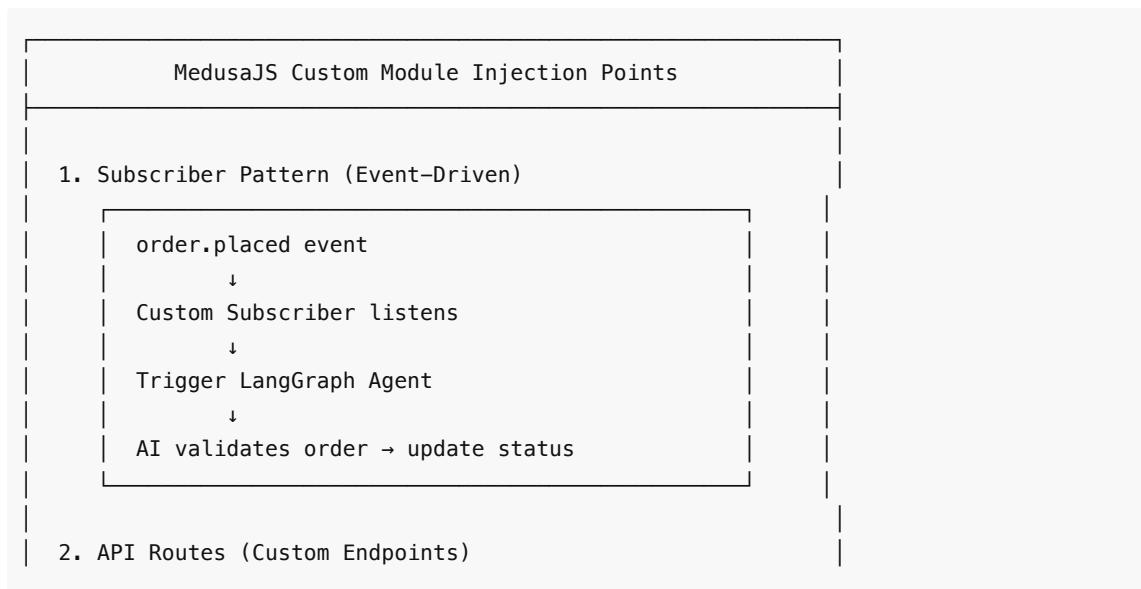
MedusaJS 2.0 B2B Architecture

Core B2B Features (Out-of-the-Box)



Custom Module Integration Points

MedusaJS allows you to inject custom logic at key lifecycle points:



```
POST /admin/sku-normalize  
POST /store/ai-cart-suggest  
POST /admin/invoice-match  
GET /store/waste-insights
```

3. Middleware (Request/Response Interception)

```
Inject AI pricing recommendations  
Add explainability metadata to cart  
Log all financial actions to audit trail
```

4. Services (Business Logic Layer)

```
class SkuNormalizationService {  
    async normalize(catalog) { ... }  
}  
class AiSuggestedCartService {  
    async generate(branch, parLevels) { ... }  
}
```

5. Workflows (Multi-Step Transactions)

```
createOrderWithGrnWorkflow:  
    step1: validate cart  
    step2: create order  
    step3: reserve inventory  
    step4: notify supplier  
    step5: schedule GRN reminder
```

LangGraph Agentic Workflows

LangGraph enables **stateful, multi-step AI reasoning** with tool-calling and human-in-the-loop approval.

LangGraph Architecture

```
LangGraph State Graph  
  
State: {  
    branch_id: "branch-001",  
    low_stock_items: [...],  
    suggested_cart: null,  
    approval_status: "pending",  
    tool_logs: []
```

```
}
```

```
Node: ProcurementAgent
- Input: low_stock_items
- Action: Call inventory API
- Output: {items, par_levels, run_rates}
```



```
Node: SkuNormalizationAgent
- Input: raw SKU names
- Action: Query vector DB for equivalents
- Output: normalized_skus
```



```
Node: SourcingAgent
- Input: normalized_skus
- Action: Compare supplier prices
- Output: {supplier, price_per_kg, lead_time}
```



```
Node: CartDraftAgent
- Input: sourcing recommendations
- Action: Calculate quantities
- Output: suggested_cart (Pydantic validated)
```

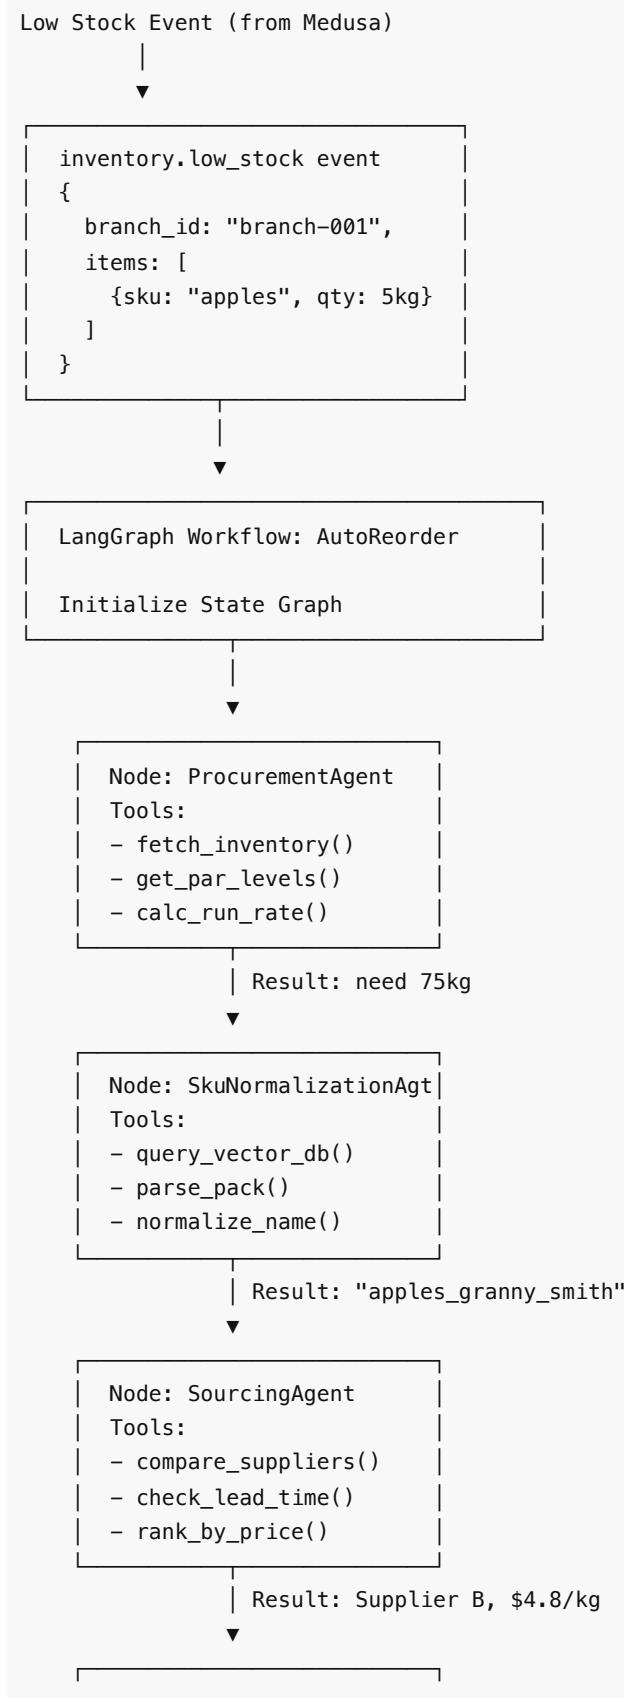


```
Node: HumanApproval (Interrupt)
- State: approval_status = "pending"
- Wait for manager input
- If approved: proceed to PO creation
- If rejected: log & exit
```



```
Node: CreatePOAgent
- Input: approved cart
- Action: Call Medusa API to create order
- Output: PO number, status
```

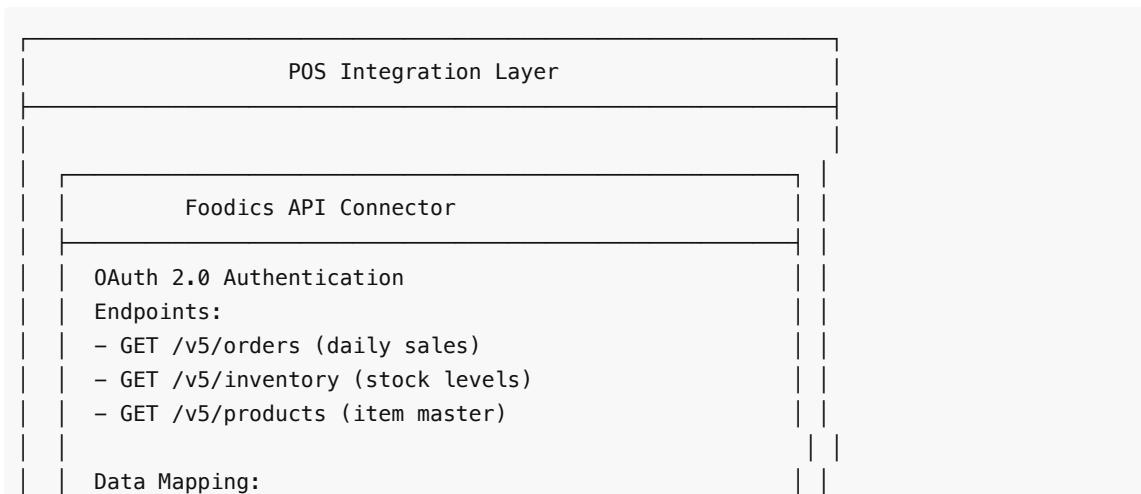
LangGraph Event Flow

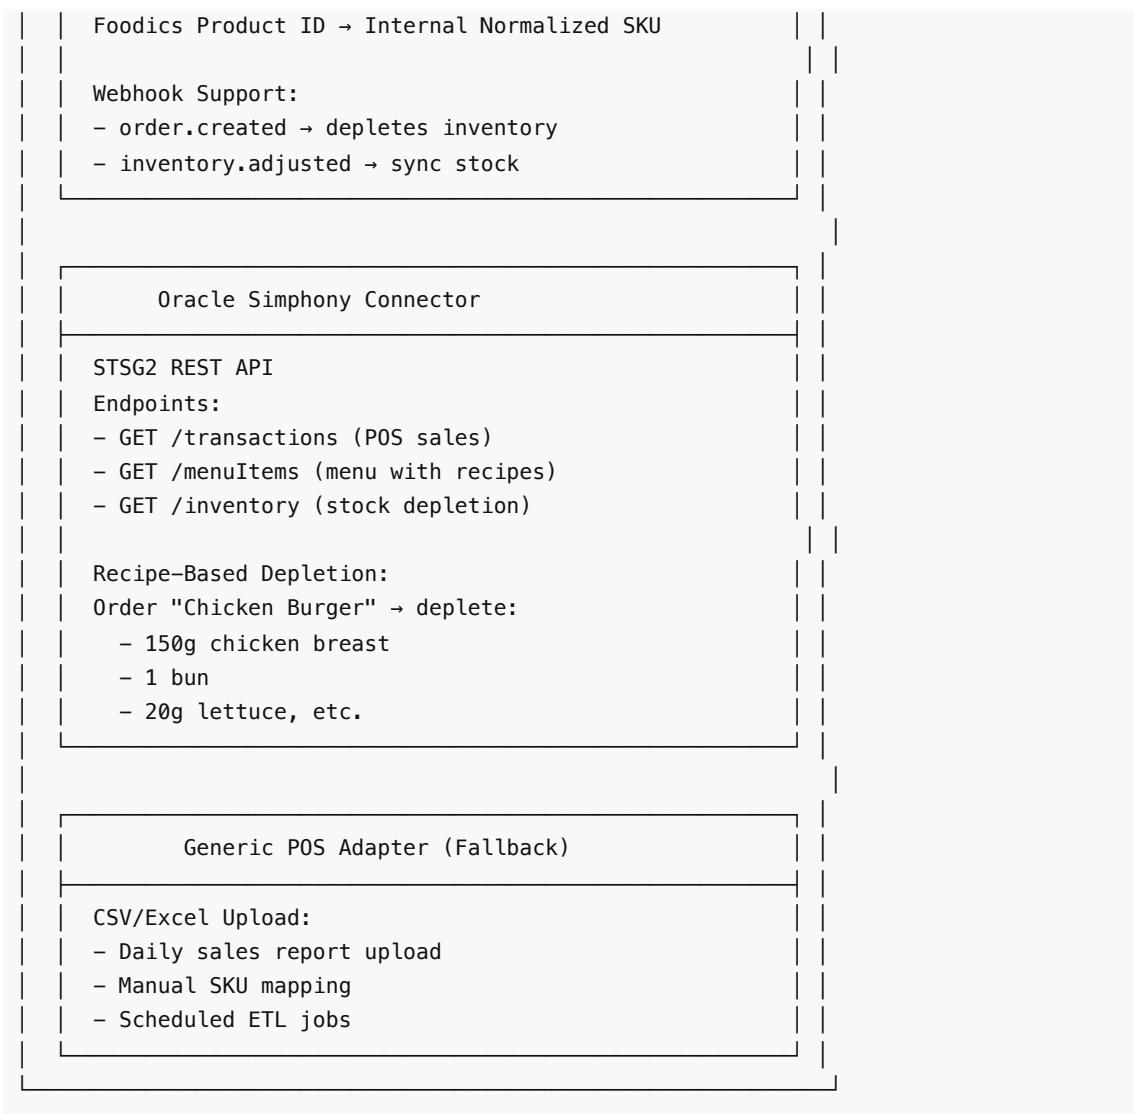




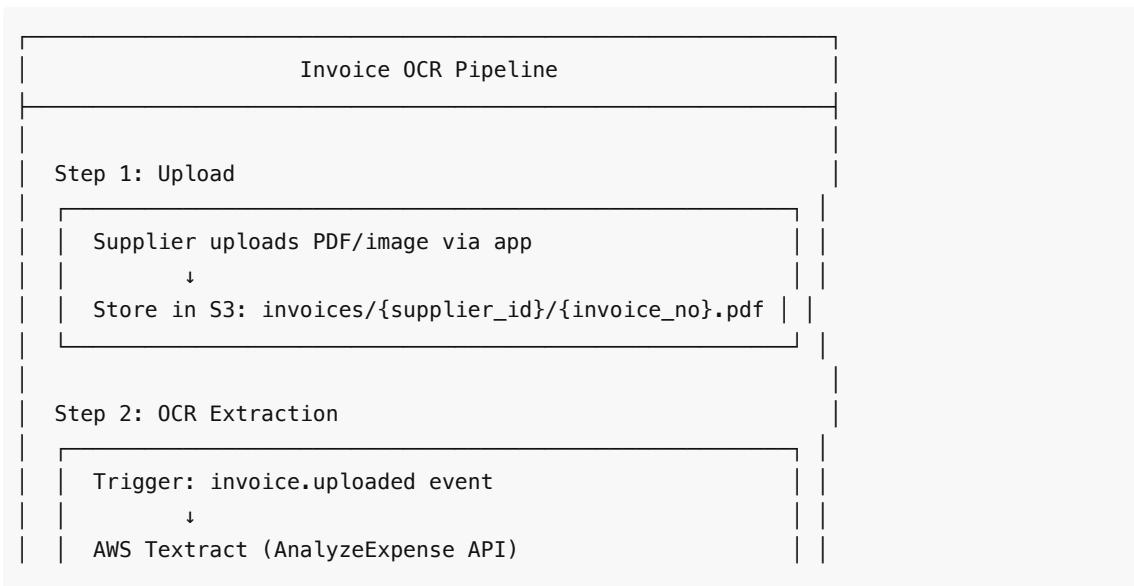
Data Acquisition Layer

POS Integration Architecture





Invoice OCR Pipeline



```
OR
Google Document AI (Invoice Parser)
    ↓
Extracted:
{
    invoice_no: "INV-12345",
    date: "2026-02-04",
    supplier: "Global Foods",
    line_items: [
        {desc: "Apples Granny Smith",
         qty: "50kg", price: "$240"},

        ...
    ],
    total: "$1,250"
}
```

Step 3: AI Validation & Normalization

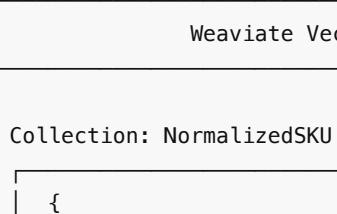
```
LangGraph Agent: InvoiceValidationAgent
- Parse line items with LLM
- Map to Normalized SKUs
- Extract pack info (50kg → 50 × 1kg)
- Calculate price per kg
```

Step 4: 2-Way / 3-Way Match

```
Compare:
- PO line items
- GRN received quantities
- Invoice billed quantities
    ↓
Flag exceptions:
- Short delivery (GRN < Invoice)
- Overcharge (Invoice price > PO price)
- Missing items
```

SKU Normalization Engine (The Moat)

Vector DB Architecture (Weaviate)



```

| id: "uuid",
|   name: "Apples Granny Smith",
|   category: "Fresh Produce",
|   attributes: {
|     grade: "A",
|     origin: "UAE",
|     organic: false
|   },
|   unit: "kilogram",
|   aliases: ["Green Apples", "Pommes Granny Smith"],
|   embedding: [0.123, 0.456, ...], ← OpenAI Ada-002
|   equiv_group: "apples_granny_smith_001"
|
}

```

Collection: SupplierCatalog

```

{
  id: "uuid",
  supplier_id: "supplier-001",
  sku: "GFD-001",
  name: "Fresh Apples Granny Smith",
  pack: "10 x 1kg",
  unit_price: 38.00,
  currency: "USD",
  normalized_sku_id: "uuid", ← Link to NormalizedSKU
  embedding: [0.120, 0.452, ...],
  last_updated: "2026-02-04"
}

```

Normalization Workflow

New Supplier Catalog Upload



Admin uploads CSV/Excel
Columns:
- SKU, Name, Pack, Price



LangGraph Agent: CatalogNormalizationAgent

For each row:
Step 1: Generate embedding
- Input: "Fresh Apples Granny Smith 10x1kg"
- Call: OpenAI embeddings API

```
- Output: [0.123, 0.456, ...]

Step 2: Similarity search in Weaviate
- Query vector DB with embedding
- Return top 3 matches with cosine similarity
- Threshold: > 0.85 = "same product"

Step 3: LLM attribute extraction
- Prompt:
  "Extract: weight, grade, origin, organic
   from 'Fresh Apples Granny Smith 10x1kg'"
- Output: {weight: 10kg, grade: null, ...}

Step 4: Pack parsing
- Input: "10 x 1kg"
- Parse: {count: 10, size: 1, unit: "kg"}
- Total weight: 10kg

Step 5: Price normalization
- unit_price / total_weight = price_per_kg
- $38 / 10kg = $3.80/kg

Step 6: Link or create Normalized SKU
- If match found: link supplier SKU
- If no match: create new Normalized SKU
- Update Weaviate
```



```
Admin Review Screen
- Show suggested matches
- Allow manual override
- Approve normalization
```

Deployment Roadmap (Step-by-Step)

Phase 0: Foundation (Weeks 1-2)

Infrastructure Setup

- ✓ Provision AWS/GCP infrastructure
- ✓ Set up PostgreSQL (RDS or managed)
- ✓ Set up Redis (ElastiCache)
- ✓ Set up S3 buckets (invoices, images, backups)
- ✓ Install Weaviate (managed or self-hosted)
- ✓ Set up CI/CD pipeline (GitHub Actions)

MedusaJS 2.0 Installation

- ✓ Initialize Medusa project
- ✓ Install B2B plugin
- ✓ Configure PostgreSQL connection
- ✓ Run migrations
- ✓ Seed initial data (regions, currencies)
- ✓ Deploy admin dashboard

Phase 1: SKU Normalization Module (Weeks 3-4)

Build Normalization Engine

- ✓ Create custom Medusa module: SkuNormalizationService
- ✓ Integrate Weaviate client
- ✓ Build CSV upload endpoint
- ✓ Implement embedding generation (OpenAI)
- ✓ Build similarity search logic
- ✓ Implement pack parser & unit converter
- ✓ Create admin UI for review & approval
- ✓ Test with 1,000 real supplier SKUs

Acceptance:

- Upload 3 supplier catalogs (300+ SKUs each)
- Achieve 90%+ auto-match accuracy
- Manual review reduces duplicates by 80%

Phase 2: POS Integration (Weeks 5-6)

Foodics Integration (Primary)

- ✓ Register OAuth app with Foodics
- ✓ Build auth flow (3-legged OAuth)
- ✓ Create FoodicsConnector service
- ✓ Implement daily sales sync (cron job)
- ✓ Map Foodics product IDs → Normalized SKUs
- ✓ Calculate theoretical inventory depletion
- ✓ Set up webhook listeners (order.created)
- ✓ Pilot with 1 restaurant (1 week)

Acceptance:

- Real-time depletion within 5 minutes of sale
- 95%+ SKU mapping accuracy
- Low stock alerts trigger within 1 hour

Phase 3: LangGraph AI Agents (Weeks 7-9)

Kitchen Copilot Agent (MVP)

- ✓ Install LangGraph + LangChain
- ✓ Build state graph: PrepPlanAgent
- ✓ Tools:
 - fetch_sales_forecast()
 - get_current_inventory()
 - get_recipes()
- ✓ Generate daily prep list:
"Based on forecast 120 orders, prep:
 - 18kg chicken breast
 - 50 burger buns
 - 3kg lettuce"
- ✓ Deploy as Medusa subscriber
- ✓ Pilot with 1 chef (2 weeks)

Acceptance:

- Prep list generated by 6am daily
- Chef feedback: 80%+ accuracy
- Reduces prep waste by 15%

Phase 4: AI Suggested Cart (Weeks 10-12)

Autopilot Purchasing Agent

- ✓ Build LangGraph workflow: AutoReorderWorkflow
- ✓ Nodes:
 - ProcurementAgent
 - SkuNormalizationAgent
 - SourcingAgent
 - CartDraftAgent
 - HumanApproval (interrupt)
 - CreatePOAgent
- ✓ Trigger: inventory.low_stock event
- ✓ Build approval UI in restaurant app
- ✓ Implement explainability (why this qty/supplier)
- ✓ Capture manager edits as feedback
- ✓ Pilot with 3 branches (4 weeks)

Acceptance:

- 70%+ of AI carts approved without edits
- 30% reduction in ordering time
- Measurable cost savings (5-10% lower \$/kg)

Phase 5: Invoice OCR & Matching (Weeks 13-15)

OCR Pipeline + Invoice Matching

- ✓ Set up AWS Textract (or Google Document AI)
- ✓ Build invoice upload endpoint
- ✓ Implement OCR extraction logic
- ✓ Build InvoiceValidationAgent (LangGraph)
- ✓ Implement 2-way match (PO vs Invoice)
- ✓ Implement 3-way match (PO vs GRN vs Invoice)
- ✓ Build exception handling UI
- ✓ Pilot with 100 invoices

Acceptance:

- OCR accuracy > 90% for structured invoices
- 2-way match reduces approval time by 60%
- 3-way match catches 95%+ of discrepancies

Phase 6: Scale & Optimize (Weeks 16-20)

Production Optimization

- ✓ Add Oracle Simphony connector
- ✓ Implement waste & variance tracking
- ✓ Build forecasting models (time series)
- ✓ Add payment/credit integrations
- ✓ Launch overseas supplier portal
- ✓ Scale to 50 restaurants
- ✓ Onboard 200 suppliers
- ✓ Process 10,000 orders/month

Tech Stack Summary

Layer	Technology
Commerce Core	MedusaJS 2.0 (Node.js, TypeScript)
Database	PostgreSQL 15+
Cache/Queue	Redis, BullMQ
Vector DB	Weaviate (or Pinecone)
AI Orchestration	LangGraph + LangChain
LLM	OpenAI GPT-4 (or Claude)

Embeddings	OpenAI Ada-002
OCR	AWS Textract or Google Document AI
POS Connectors	Foodics API, Oracle Symphony STSG2
Frontend	Next.js 14+ (React, TypeScript)
Storage	AWS S3 (or GCP Cloud Storage)
Hosting	AWS ECS/EKS or GCP Cloud Run
CI/CD	GitHub Actions
Monitoring	Datadog or New Relic

Key Benefits of This Architecture

- Modular:** Medusa's plugin system keeps AI logic decoupled from core commerce
- Event-Driven:** Low stock triggers AI without polling
- Stateful AI:** LangGraph maintains context across multi-step workflows
- Human-in-the-Loop:** Manager approval is a first-class citizen
- Explainable:** Every AI decision includes reasoning
- Scalable:** PostgreSQL + Redis + Weaviate handle 1M+ SKUs
- UAE-Optimized:** Multi-currency, Arabic support, local POS integrations
- B2B Native:** Company accounts, RFQs, credit terms out-of-the-box